

CEC 327

Assignment 3: Distributed File System

Professor Oscar Morales

Due date April 5th

1 Objective

The objectives of this project are:

- 1) Identify the advantages RPC for programming distributed applications.
- 2) Understand the use of middlewares (in particular Java RMI) to implement distributed operative systems
- 3) Identify the advantages of distributed file system

The task is to develop a distributed file system based on Chord (a peer-to-peer system.) In Chord, the processes and the files are mapped to a unique identifier using a hash function that distributes them uniformly. Thus, when a node joins or leaves the systems, only $O(\log n)$ keys are needed to be move where n is the number of processes. Each process stores the IP and port of the peers that have joined.

In Chord, process and files are hashed to an m -bit unique identifier where m is sufficiently large so that the probability of collisions during the hash is negligible, for example using an MD5 digest. The Chord overlay uses an ordered logical ring of size 2^m . A key k gets assigned to the first node such that its node identifier equals or follows the key identifier of k in the common identifier space. Each process i maintains a routing table, called the finger table, with at most $O(\log n)$ distinct entries, such that the x -th entry ($1 \leq x \leq m$) is the process of the node $\text{succ}(i + 2^{x-1})$. The following code uses an RPC notation.

(variables) integer: $successor \leftarrow initial_value$;
 integer: $predecessor \leftarrow initial_value$;
 integer $finger[1...m]$;
 integer: $next_finger \leftarrow 1$;
 (1) $i.Locate_Successor(key)$; where $key \neq i$:
 (1a) if $key \in (i, successor]$ then
 (1b) return($successor$)
 (1c) else
 (1d) $j \leftarrow Closest_Preceding_Node(key)$;
 (1e) return ($j.Locate_Successor(key)$).
 (2) $i.Closest_Preceding_Node(key)$; where $key \neq i$:
 (2a) for $count = m$ down to 1 do
 (2b) if $finger[count] \in (i, key]$ then
 (2c) break();
 (2d) return($finger[count]$).
 (3) $i.Create_New_Ring$:
 (3a) $predecessor \leftarrow \perp$;
 (3b) $successor \leftarrow i$.
 (4) $i.Join_Ring(j)$, where j is any node on the ring to be joined:
 (4a) $predecessor \leftarrow \perp$;
 (4b) $successor \leftarrow j.Locate_Successor(i)$.
 (5) $i.Stabilize$: executed periodically to verify and inform successor
 (5a) $x \leftarrow successor.predecessor$;
 (5b) if $x \in (i, successor)$ then
 (5c) $successor \leftarrow x$;
 (5d) $successor.Notify(i)$.
 (6) $i.Notify(j)$: j believes it is predecessor of i
 (6a) if $predecessor = \perp$ or $j \in (predecessor, i)$ then
 (6b) transfer keys not in the range $(j, i]$ to j ;
 (6c) $predecessor \leftarrow j$.
 (7) $i.Fix_Fingers$: executed periodically to update the finger table
 (7a) $next_finger \leftarrow next_finger + 1$;
 (7b) if $next_finger > m$ then
 (7c) $next_finger \leftarrow 1$;
 (7d) $finger[next_finger] \leftarrow Locate_Successor(i + 2^{next_finger-1})$.
 (8) $i.Check_Predecessor$: executed periodically to verify whether predecessor still exists
 (8a) if $predecessor$ has failed then
 (8b) $predecessor \leftarrow \perp$.

2 Basic P2P Programming Interface

- *Join(ip, port)*: Join the system. First, peers obtain the GUID base on the IP and Port.
- *put(GUID, data)*: Stores data in the processor responsible for storing the object.
- *value = get(GUID)*: Retrieves the data associated with GUID from one of the nodes responsible for it.
- *remove(GUID)*: Deletes all references to GUID and the associated data.
- *print*: Print the state of the systems.

The GUID is the file name is the MD5 of the file module 2^{31} , i.e., $GUID_{fileName} = MD5(fileName) \bmod 2^{31}$

3 Distributed File System

To store big files, the file will be split into several pages where each page is a full file. A metadata will store all the information of the files. We will use JSON to maintain all the files information. For example:

```
{
  "metadata" :
  {
    file :
    {
      name : "File1"
      numberOfPages : "3"
      pageSize : "1024"
      size : "2291"
      page :
      {
        number : "1"
        guid : "22412"
```

```

        size    : "1024"
    }
    page :
    {
        number : "2"
        guid   : "46312"
        size   : "1024"
    }
    page :
    {
        number : "3"
        guid   : "93719"
        size   : "243"
    }
}
}
}

```

The metadata will be store in the peer-to-peer with guid md5("Metadata") Whenever a user wants to access a file, the DFS will read and parse the meta-data file. If the operation changes the metadata content, then the file will be updated. The DFS must support: join (join the DFS), ls (list all the files), touch (creates a file), delete, read (read a given page of the file), tail (read the last page of the file), head (read the first page of the file), append (add content at the end of the file), move (rename the file).

4 Grading

Criteria	Weight
Documentation of your program	20%
Execution	80%

The documentation must be generated using Doxygen or Javadocs.

5 Deliverables

Compress the source and documentation in a zip file and upload to beachboard. The zipfile must contain two folders:

- 1.- Src: All the source code to compile it
- 2.- Docs: HTML with the documentation. It has to contain the definitions of the methods with a short description, parameters and output of the method.